



# RISC-V Capacity and Bandwidth QoS Register Interface

RISC-V CMQRI Task Group

Version 0.1, 04/2023: This document is in development. Assume everything can change. See  
<http://riscv.org/spec-state> for details.

# Table of Contents

Preamble.....	1
Copyright and license information.....	2
Contributors.....	3
1. Introduction.....	4
2. QoS Identifiers.....	6
2.1. Associating <b>RCID</b> and <b>MCID</b> with requests.....	6
2.1.1. RISC-V hart initiated requests (Ssqosid).....	6
2.1.2. Device initiated requests.....	7
2.2. Access-type ( <b>AT</b> ).....	7
3. Capacity-controller QoS Register Interface.....	9
3.1. Capacity-controller capabilities ( <b>cc_capabilities</b> ).....	9
3.2. Capacity usage monitoring control ( <b>cc_mon_ctl</b> ).....	10
3.3. Capacity usage monitoring counter value ( <b>cc_mon_ctr_val</b> ).....	12
3.4. Capacity allocation control ( <b>cc_alloc_ctl</b> ).....	13
3.5. Capacity block mask ( <b>cc_block_mask</b> ).....	15
4. Bandwidth-controller QoS Register Interface.....	17
4.1. Capabilities ( <b>bc_capabilities</b> ).....	17
4.2. Bandwidth usage monitoring control ( <b>bc_mon_ctl</b> ).....	18
4.3. Bandwidth monitoring counter value ( <b>bc_mon_ctr_val</b> ).....	20
4.4. Bandwidth allocation control ( <b>bc_alloc_ctl</b> ).....	21
4.5. Bandwidth allocation configuration ( <b>bc_bw_alloc</b> ).....	22
5. IOMMU extension for QoS ID.....	25
5.1. IOMMU registers.....	25
5.1.1. Reset behavior.....	25
5.1.2. IOMMU capabilities ( <b>capabilities</b> ).....	25
5.1.3. IOMMU QoS ID ( <b>iommu_qosid</b> ).....	26
5.2. Device-context fields.....	27
5.3. IOMMU ATC capacity allocation and monitoring.....	27
6. Hardware Guidelines.....	28
6.1. Sizing QoS Identifiers.....	28
6.2. Sizing monitoring counters.....	28
7. Software Guidelines.....	29
7.1. Reporting capacity and bandwidth controllers.....	29
7.2. Context switching QoS Identifiers.....	29
7.3. QoS configurations for virtual machines.....	29
7.4. QoS Identifiers for supervisor and machine mode.....	30
Bibliography.....	31

# Preamble



*This document is in the [Development state](#)*

Assume everything can change. This draft specification will change before being accepted as standard, so implementations made to this draft specification will likely not conform to the future standard.

# Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at [creativecommons.org/licenses/by/4.0/](https://creativecommons.org/licenses/by/4.0/).

Copyright 2022 by RISC-V International.

# Contributors

This RISC-V specification has been contributed to directly or indirectly by:

Aaron Durbin, Adrien Ricciardi, Allen Baum, Allison Randal, Ambika Krishnamoorthy, Daniel Gracia Pérez, David Kruckemeyer, David Weaver, Derek Hower, Drew Fustini, Eric Shiu, Greg Favor, Ravi Sahita, Vedvyas Shanbhogue

# Chapter 1. Introduction

Quality of Service (QoS) is defined as the minimal end-to-end performance that is guaranteed in advance by a service level agreement (SLA) to a workload. A workload may be a single application, a group of application, a virtual machine, a group of virtual machines, or a combination of those. The performance may be measured in the form of metrics such as instructions per cycle (IPC), latency of servicing work, etc.

Various factors such as the available cache capacity, memory bandwidth, interconnect bandwidth, CPU cycles, system memory, etc. affect the performance in a computing system that runs multiple workloads concurrently. Furthermore, when there is arbitration for shared resources, the prioritization of the workloads requests against other competing requests may also affect the performance of the workload. Such interference due to resource sharing may lead to unpredictable workload performance [1].

When multiple workloads are running concurrently on modern processors with large core counts, multiple cache hierarchies, and multiple memory controllers, the performance of a workload becomes less deterministic or even non-deterministic. This is because the performance depends on the behavior of all the other workloads in the machine that contend for the shared resources leading to interference. In many deployment scenarios such as public cloud servers the workload owner may not be in control of the type and placement of other workloads in the platform.

System software can control some of these resources available to the workload, such as the number of hardware threads made available for execution, the amount of system memory allocated to the workload, the number of CPU cycles provided for execution, etc.

System software needs additional tools to control interference to a workload and thereby reduce the variability in performance experienced by one workload due to other workload's cache capacity usage, memory bandwidth usage, interconnect bandwidth usage, etc. through a resource allocation capability. The resource allocation capability enables system software to reserve capacity and/or bandwidth to meet the performance goals of the workload. Such controls enable improving the utilization of the system by collocating workloads while minimizing the interference caused by one workload to another [2].

Effective use of the resource allocation capability requires hardware to provide a resource monitoring capability by which the resource requirements of a workload needed to meet a certain performance goal can be characterized. A typical use model involves profiling the resource usage of the workload using the resource monitoring capability and to establish resource allocations for the workload using the resource allocation capability.

The allocation may be in the form of capacity or bandwidth depending on the type of resource. For caches, TLBs, and directories, the resource allocation is in the form of storage capacity. For interconnects and memory controllers the resource allocation is in the form of bandwidth.

Workloads generate different types of accesses to shared resources. For example, some of the requests may be for accessing instructions and others may be to access data operated on by the workload. Certain data accesses may not have temporal locality whereas others may have a high probability of reuse. In some cases it is desirable to provide differentiated treatment to each type of access by providing unique resource allocation to each access type.

RISC-V Capacity and Bandwidth Controller QoS Register Interface (CBQRI) specification specifies:

1. QoS identifiers to identify workloads that originate requests to the shared resources. These QoS identifiers include an identifier for resource allocation configurations and an identifier for the monitoring counters used to monitor resource usage. These identifiers accompany each request made by the workload to the shared resource. [Chapter 2](#) specifies the mechanism to associate the identifiers with workloads.
2. Access-type identifiers to accompany request to access a shared resource to allow differentiated treatment of each access-type (e.g., code vs. data, etc.). The access-types are defined in [Chapter 2](#).
3. Register interface for capacity allocation in controllers such as shared caches, directories, etc. The capacity allocation register interface is specified in [Chapter 3](#).
4. Register interface for capacity usage monitoring. The capacity usage monitoring register interface is specified in [Chapter 3](#).
5. Register interface for bandwidth allocation in controllers such as interconnect and memory controllers. The bandwidth allocation register interface is specified in [Chapter 4](#).
6. Register interface for bandwidth usage monitoring. The bandwidth usage monitoring register interface is specified in [Chapter 4](#).

The capacity and bandwidth controller register interfaces for resource allocation and usage monitoring are defined as memory-mapped registers. Each controller that supports CBQRI provides a set of registers that are located at a range of physical address space that is a multiple of 4-KiB and the lowest address of the range is aligned to 4-KiB. The memory-mapped registers may be accessed using naturally aligned 4-byte or 8-byte memory accesses. The controller behavior for register accesses where the address is not aligned to the size of the access, or if the access spans multiple registers, or if the size of the access is not 4-bytes or 8-bytes, is **UNSPECIFIED**. The atomicity of access to an 8-byte register is **UNSPECIFIED**. A 4-byte access to a register must be single-copy atomic.

The controller registers have little-endian byte order (even for systems where all harts are big-endian-only).



Big-endian-configured harts that make use of the register interface may implement the **REV8** byte-reversal instruction defined by the Zbb extension. If **REV8** is not implemented, then endianness conversion may be implemented using a sequence of instructions.

A controller may support a subset of capabilities defined by CBQRI. When a capability is not supported the registers and/or fields used to configure and/or control such capabilities are hardwired to 0. Each controller supports a capabilities register to enumerate the supported capabilities.

# Chapter 2. QoS Identifiers

Monitoring or allocation of resources requires a way to identify the originator of the request to access the resource.

CBQRI and the Ssqosid extension provides a mechanism by which a workload can be associated with a resource control ID (**RCID**) and a monitoring counter ID (**MCID**) that accompany each request made by the workload to shared resources.

To provide differentiated services to workloads, CBQRI defines a mechanism to configure resource usage limits, in the form of capacity or bandwidth, per supported access type, for an **RCID** in the resource controllers that control accesses to such shared resources.

To monitor the resource utilization by a workload CBQRI defines a mechanism to configure counters identified by the **MCID** to count events in the resource controllers that control accesses to such shared resources.

## 2.1. Associating **RCID** and **MCID** with requests

### 2.1.1. RISC-V hart initiated requests (Ssqosid)

The Ssqosid extension, introduces a read/write S/HS-mode register (**sqoscfg**) to configure QoS Identifiers to be used with requests made by the hart to shared resources.

The **sqoscfg** CSR is a 32-bit S/HS-mode read/write WARL register to configure a resource control ID (**RCID**) and a monitoring counter ID (**MCID**). The **RCID** and **MCID** accompany each request made by the hart to shared resources such as interconnects, caches, memory, etc.




Figure 1. **sqoscfg** register for RV32 and RV64



The type of request made to the shared resource controller depends on the type of shared resource. In case of resources such as caches or memory these may be a memory access request. In case of resources such as CPU dispatch slots and retirement bandwidth the request may be to allocate such resources for execution.

The **RCID** in the request is used by the resource controllers to determine the resource allocations (e.g., cache occupancy limits, memory bandwidth limits, etc.) to enforce. The **MCID** in the request is used by the resource controllers to identify the ID of a counter to monitor resource usage (e.g., cache occupancy, memory bandwidth, etc.).

Access to **sqoscfg** when **V=1** causes a virtual-instruction exception.



At reset it is suggested that the **RCID** field of **sqoscfg** be set to 0 as typically the resource controllers in the SoC default to a reset behavior of associating all capacity or bandwidth to the **RCID** value of 0.



The value of **MCID** at reset, unlike the **RCID**, does not affect functional behavior. Implementations may choose a convenient legal value for the **MCID** reset value.

The **RCID** and **MCID** configured in the **sqoscfg** CSR apply to all privilege modes of software execution on that hart.

The **QOSE** (bit 58) field of the **menvcfg** (and **menvcfgh**) machine-level CSR controls if the **sqoscfg** CSR is accessible in modes less privileged than M. When **QOSE** is 0, attempts to access **sqoscfg** CSR at privilege modes less privileged than M cause an illegal instruction exception.

### 2.1.2. Device initiated requests

Devices may be configured with an **RCID** and **MCID** for requests originated from the device if the device implementation supports such capability. The method to configure the QoS identifiers into devices is **UNSPECIFIED**.

Where the device does not natively support being configured with an **RCID** and **MCID**, the implementation may provide a shim at the device interface that may be configured with the **RCID** and **MCID** that are associated with requests originating from the device. The method to configure such QoS identifiers into a shim is **UNSPECIFIED**.

If the system supports an IOMMU, then the IOMMU may be configured with the **RCID** and **MCID** to associate requests from the device with QoS identifiers. The RISC-V IOMMU [3] extension to support configuring QoS identifiers is specified in [Chapter 5](#).

## 2.2. Access-type (**AT**)

In some usages, in addition to providing differentiated service among workloads, the ability to differentiate between resource usage for accesses made by the same workload may be required. For example, the capacity allocated in a shared cache for code storage may be differentiated from the capacity allocated for data storage and thereby avoid code from being evicted from such shared cache due to a data access.

When differentiation based on access type (e.g. code vs. data) is supported the requests also carry an access-type (**AT**) indicator. The resource controllers may be configured with separate capacity and/or bandwidth allocations for each supported access-type. CBQRI defines a 3-bit **AT** field, encoded as specified in [Table 1](#), in the register interface to configure differentiated resource allocation and monitoring for each **AT**.

Table 1. Encodings of **AT** field

Value	Name	Description
0	<b>Data</b>	Requests to access data.
1	<b>Code</b>	Requests for code execution.
2-5	Reserved	Reserved for future standard use.
6-7	Custom	Designated for custom use.

If a request is received with an unsupported **AT** value then the resource controller behaves as if the

AT value was 0.

# Chapter 3. Capacity-controller QoS Register Interface

Controllers, such as cache controllers, that support capacity allocation and usage monitoring provide a memory-mapped capacity-controller QoS register interface.

Table 2. Capacity-controller QoS register layout

Offset	Name	Size	Description	Optional?
0	<code>cc_capabilities</code>	8	Capabilities	No
8	<code>cc_mon_ctl</code>	8	Usage monitoring control	Yes
16	<code>cc_mon_ctr_val</code>	8	Monitoring counter value	Yes
24	<code>cc_alloc_ctl</code>	8	Capacity allocation control	Yes
32	<code>cc_block_mask</code>	M * 8	Capacity block mask	Yes

The size and offset in Table 2 are specified in bytes.

The size of the `cc_block_mask` register is determined by the `NCBLKS` field of the `cc_capabilities` register but is always a multiple of 8 bytes.

The reset value is 0 for the following registers fields.

- `cc_mon_ctl` - `BUSY` field
- `cc_alloc_ctl` - `BUSY` field

The reset value is `UNSPECIFIED` for all other registers and/or fields.

The capacity controllers at reset must allocate all available capacity to `RCID` value of 0. When the capacity controller supports capacity allocation per access-type, then all available capacity is shared by all the access-type for `RCID=0`. The capacity allocation for all other `RCID` values is `UNSPECIFIED`. The capacity controller behavior for handling a request with a non-zero `RCID` value before configuring the capacity controller with capacity allocation for that `RCID` is `UNSPECIFIED`.

## 3.1. Capacity-controller capabilities (`cc_capabilities`)

The `cc_capabilities` register is a read-only register that holds the capacity-controller capabilities.




Figure 2. Capacity-controller capabilities register fields

The **VER** field holds the version of the specification implemented by the capacity controller. The low nibble is used to hold the minor version of the specification and the upper nibble is used to hold the major version of the specification. For example, an implementation that supports version 1.0 of the specification reports 0x10.

The **NCBLKS** field holds the total number of allocatable capacity blocks in the controller. The capacity represented by an allocatable capacity block is **UNSPECIFIED**. The capacity controllers support allocating capacity in multiples of an allocatable capacity block.



For example, a cache controller that supports capacity allocation by ways may report the number of ways as the number of allocatable capacity blocks.

If **FRCID** is 1, the controller supports an operation to flush and deallocate the capacity blocks occupied by an **RCID**.

## 3.2. Capacity usage monitoring control (**cc\_mon\_ctl**)

The **cc\_mon\_ctl** register is used to control monitoring of capacity usage by a **MCID**. When the controller does not support capacity usage monitoring the **cc\_mon\_ctl** register is read-only zero.




Figure 3. Capacity usage monitoring control (**cc\_mon\_ctl**)

Capacity controllers that support capacity usage monitoring implement a usage monitoring counter for each supported **MCID**. The usage monitoring counter may be configured to count a monitoring event. When an event matching the event configured for the **MCID** occurs then the monitoring counter is updated. The event matching may optionally be filtered by the access-type.

The **OP**, **AT**, **ATV**, **MCID**, and **EVT\_ID** fields of the register are WARL fields.

The **OP** field is used to instruct the controller to perform an operation listed in Table 3. The supported operations of the **OP** field include configuring an event to count in the monitoring counter and obtaining a snapshot of the counter value.

The **EVT\_ID** field is used to program the identifier of the event to count in the monitoring counter selected by **MCID**. The **AT** field is used to program the access-type to count, and its validity is indicated by the **ATV** field. When **ATV** is 0, the counter counts requests with all access-types, and the **AT** value is ignored.

When the **EVT\_ID** for a **MCID** is programmed with a non-zero and legal value, the counter is reset to 0 and starts counting matching events for requests with the matching **MCID** and **AT** (if **ATV** is 1). However, if the **EVT\_ID** is configured as 0, the counter retains its current value but stops counting.

A controller that does not support monitoring by access-type can hardwire the **ATV** and the **AT** fields to 0, indicating that the counter counts requests with all access-types.

Table 3. Capacity usage monitoring operations (**OP**)

Operation	Encoding	Description
—	0	Reserved for future standard use.
CONFIG_EVENT	1	Configure the counter selected by MCID to count the event selected by EVT_ID, AT, and ATV. The EVT_ID encodings are listed in Table 4.
READ_COUNTER	2	Snapshot the value of the counter selected by MCID into cc_mon_ctr_val register. The EVT_ID, AT, and ATV fields are not used by this operation.
—	3-23	Reserved for future standard use.
—	24-31	Designated for custom use.

Table 4. Capacity usage monitoring event ID (EVT\_ID)

Event ID	Encoding	Description
None	0	Counter does not count and retains its value.
Occupancy	1	Counter is incremented by 1 when a request with a matching MCID and AT allocates a unit of capacity. The counter is decremented by 1 when a unit of capacity is de-allocated.
—	2-127	Reserved for future standard use.
—	128-256	Designated for custom use.

When the `cc_mon_ctl` register is written, the controller may need to perform several actions that may not complete synchronously with the write. A write to the `cc_mon_ctl` sets the read-only `BUSY` bit to 1 indicating the controller is performing the requested operation. When the `BUSY` bit reads 0, the operation is complete, and the read-only `STATUS` field provides a status value (see Table 5 for details). Written values to the `BUSY` and the `STATUS` fields are ignored. An implementation that can complete the operation synchronously with the write may hardwire the `BUSY` bit to 0. The state of the `BUSY` bit, when not hardwired to 0, shall only change in response to a write to the register. The `STATUS` field remains valid until a subsequent write to the `cc_mon_ctl` register.

Table 5. `cc_mon_ctl.STATUS` field encodings

STATUS	Description
0	Reserved
1	The operation was successfully completed.
2	An invalid operation (OP) was requested.
3	An operation was requested for an invalid MCID.
4	An operation was requested for an invalid EVT_ID.
5	An operation was requested for an invalid AT.
6-63	Reserved for future standard use.
64-127	Designated for custom use.

When the `BUSY` bit is set to 1, the behavior of writes to the `cc_mon_ctl` is `UNSPECIFIED`. Some implementations may ignore the second write, while others may perform the operation determined by the second write. To ensure proper operation, software must first verify that the `BUSY` bit is 0

before writing the `cc_mon_ctl` register.

### 3.3. Capacity usage monitoring counter value (`cc_mon_ctr_val`)

The `cc_mon_ctr_val` is a read-only register that holds a snapshot of the counter selected by the `READ_COUNTER` operation. When the controller does not support capacity usage monitoring, the `cc_mon_ctr_val` register always reads as zero.




Figure 4. Capacity usage monitoring counter value (`cc_mon_ctr_val`)

The counter is valid if the `INV` field is 0. The counter may be marked `INV` if it underflows or the controller, for `UNSPECIFIED` reasons determine the count to be not valid. The counters marked `INV` may become valid in future.



A counter may underflow when capacity is de-allocated following a reset of the counter to 0. This may be due to the `MCID` being reallocated to a new workload while the capacity controller still holds capacity allocated by the workload to which the `MCID` was previously allocated. The counter value should typically stabilize to reflect the capacity usage of the new workload after the workload has executed for a short duration following the counter reset.



Some implementations may not store the `MCID` of the request that caused the capacity to be allocated with every unit of capacity in the controller to optimize on the storage overheads. Such controllers may in turn rely on statistical sampling to report the capacity usage by tagging only a subset of the capacity units.

Set-sampling is a technique commonly used in caches to estimate the cache occupancy with a relatively small sample size. The basic idea behind set-sampling is to select a subset of the cache sets and monitor only those sets. By keeping track of the hits and misses in the monitored sets, it is possible to estimate the overall cache occupancy with a high degree of accuracy. The size of the subset needed to obtain accurate estimates depends on various factors, such as the size of the cache, the cache access patterns, and the desired accuracy level. Research [4] has shown that set-sampling can provide statistically accurate estimates with a relatively small sample size, such as 10% or less, depending on the cache properties and sampling technique used.

When the controller has not observed enough samples to provide an accurate value in the monitoring counter, it may report the counter as being `INV` until more accurate measurements are available. This helps to prevent inaccurate or misleading data from being used in capacity planning or other decision-making processes.

### 3.4. Capacity allocation control (**cc\_alloc\_ctl**)

The **cc\_alloc\_ctl** register is used to configure allocation of capacity to an **RCID** per access-type (**AT**). The **RCID** and **AT** fields in this register are WARL. If a controller does not support capacity allocation then this register is read-only zero. If the controller does not support capacity allocation per access-type then the **AT** field is read-only zero.




Figure 5. Capacity allocation control (**cc\_alloc\_ctl**)

The **OP**, **AT**, and **RCID** are WARL fields.

The **OP** field used to instruct the capacity controller to perform an operation listed in Table 6. Some operations require specifying the capacity blocks to operate on. The capacity blocks, in the form of a bitmask, for such operations are specified in the **cc\_block\_mask** register. To request operations that need a capacity block mask to be specified, software must first program the **cc\_block\_mask** register and then request the operation using the **cc\_alloc\_ctl** register.

Table 6. Capacity allocation operations (**OP**)

Operation	Encoding	Description
—	0	Reserved for future standard use.
<b>CONFIG_LIMIT</b>	1	Configure a capacity allocation for requests by <b>RCID</b> and of access-type <b>AT</b> . The capacity allocation is specified in the <b>cc_block_mask</b> register.
<b>READ_LIMIT</b>	2	Read back the previously configured capacity allocation for requests by <b>RCID</b> and of access-type <b>AT</b> . The configured allocation is returned as a bit-mask in the <b>cc_block_mask</b> register on successful completion of the operation.
<b>FLUSH_RCID</b>	3	Deallocate the capacity used by the specified <b>RCID</b> and access-type <b>AT</b> . This operation is supported if the <b>capabilities.FRCID</b> bit is 1.+ The <b>cc_block_mask</b> register is not used for this operation.+ The configured capacity allocation is not changed by this operation.
—	4-23	Reserved for future standard use.
—	24-31	Designated for custom use.

Capacity controllers enumerate the allocatable capacity blocks in the **NBLKS** field of the **cc\_capabilities** register. The **cc\_block\_mask** register is programmed with a bit-mask where each bit represents a capacity block for the operation.

A capacity allocation must be configured for each supported access-type by the controller. An implementation that does not support capacity allocation per access-type may hardwire the **AT** field to 0 and associate use the same capacity allocation configuration for requests with all access-types. When capacity allocation per access-type is supported, identical limits may be configured for two or more access-types if different capacity allocation per access-type is not required. If capacity is not

allocated for each access-type supported by the controller, the behavior is **UNSPECIFIED**.



A cache controller that supports capacity allocation indicates the number of allocatable capacity blocks in **cc\_capabilities.NCBLKS** field. For example, let's consider a cache with **NCBLKS=8**. In this example, the **RCID=5** has been allocated capacity blocks numbered 0 and 1 for requests with access-type **AT=0**, and has been allocated capacity blocks numbered 2 for requests with access-type **AT=1**. The **RCID=3** in this example has been allocated capacity blocks numbered 3 and 4 for both **AT=0** and **AT=1** access-types as separate capacity allocation by access-type is not required for this workload. Further in this example, the **RCID=6** has been configured with the same capacity block allocations as **RCID=3**. This implies that they share a common capacity allocation in this cache but may have been associated with different **RCID** to allow differentiated treatment in another capacity and/or bandwidth controller.

	7	6	5	4	3	2	1	0
<b>RCID=3, AT=0</b>	0	0	0	1	1	0	0	0
<b>RCID=3, AT=1</b>	0	0	0	1	1	0	0	0
<b>RCID=5, AT=0</b>	0	0	0	0	0	0	1	1
<b>RCID=5, AT=1</b>	0	0	0	0	0	1	0	0
<b>RCID=6, AT=0</b>	0	0	0	1	1	0	0	0
<b>RCID=6, AT=1</b>	0	0	0	1	1	0	0	0

The **FLUSH\_RCID** operation may incur a long latency to complete. New requests to the controller by the **RCID** being flushed are allowed. Additionally, the controller is allowed to deallocate capacity that was allocated after the operation was initiated.



For cache controllers, the **FLUSH\_RCID** operation may perform an operation similar to that performed by the RISC-V **CBO.FLUSH** instruction on each cache block that is part of the allocation configured for the **RCID**.

The **FLUSH\_RCID** operation can be used as part of reclaiming a previously allocated **RCID** and associating it with a new workload. When such a reallocation is performed, the capacity controllers may have capacity allocated by the old workload and thus for a short warmup duration the capacity controller may be enforcing capacity allocation limits that reflect the usage by the old workload. Such warmup durations are typically not statistically significant, but if that is not desired, then the **FLUSH\_RCID** operation can be used to flush and evict capacity allocated by the old workload.

When the **cc\_alloc\_ctl** register is written, the controller may need to perform several actions that may not complete synchronously with the write. A write to the **cc\_alloc\_ctl** sets the read-only **BUSY** bit to 1 indicating the controller is performing the requested operation. When the **BUSY** bit reads 0, the operation is complete, and the read-only **STATUS** field provides a status value (Table 7) of the requested operation. Values written to the **BUSY** and the **STATUS** fields are always ignored. An



implementation that can complete the operation synchronously with the write may hardwire the **BUSY** bit to 0. The state of the **BUSY** bit, when not hardwired to 0, shall only change in response to a write to the register. The **STATUS** field remains valid until a subsequent write to the **cc\_alloc\_ctl** register.

Table 7. **cc\_alloc\_ctl.STATUS** field encodings

<b>STATUS</b>	<b>Description</b>
0	Reserved
1	The operation was successfully completed.
2	An invalid or unsupported operation ( <b>OP</b> ) requested.
3	An operation was requested for an invalid <b>RCID</b> .
4	An operation was requested for an invalid <b>AT</b> .
5	An invalid capacity block mask was specified.
6-63	Reserved for future standard use.
64-127	Designated for custom use.

When the **BUSY** bit is set to 1, the behavior of writes to the **cc\_alloc\_ctl** register or to the **cc\_block\_mask** register is **UNSPECIFIED**. Some implementations may ignore the second write and others may perform the operation determined by the second write. To ensure proper operation, software must verify that **BUSY** bit is 0 before writing the **cc\_alloc\_ctl** register or the **cc\_block\_mask** register.

### 3.5. Capacity block mask (**cc\_block\_mask**)

The **cc\_block\_mask** is a WARL register. If the controller does not support capacity allocation i.e. **NCBLKS** is 0, then this register is read-only 0.

The register has **NCBLKS** bits each corresponding to one allocatable capacity block in the controller. The width of this register is variable but always a multiple of 64 bits. The bitmap width in bits (**BMW**) is determined in [equation \(1\)](#). The division operation in this equation is an integer division.

$$BMW = \frac{NCBLKS + 63}{64} \times 64 \quad (1)$$

Bits **NCBLKS-1:0** are read-write, and the bits **BMW-1:NCBLKS** are read-only and have a value of 0.

The process of configuring capacity allocation for an **RCID** and **AT** begins by programming the **cc\_block\_mask** register with a bit-mask that identifies the capacity blocks to be allocated. Next, the **cc\_alloc\_ctl** register is written to request a **CONFIG\_LIMIT** operation for the **RCID** and **AT**. Once a capacity allocation limit has been established, a request may be allocated capacity in the capacity blocks allocated to the **RCID** and **AT** associated with the request. It is important to note that at least one capacity block must be allocated using **cc\_block\_mask** when allocating capacity, or else the operation will fail with **STATUS=5**. Overlapping capacity block masks among **RCID** and/or **AT** are allowed to be configured.



A set-associative cache controller that supports capacity allocation by ways can advertise **NCBLKS** as the number of ways per set in the cache. To Allocate capacity in such a cache for an **RCID** and **AT**, a subset of ways must be selected and mask of the selected ways must be programmed in **cc\_block\_mask** when requesting the **CONFIG\_LIMIT** operation.

To read the capacity allocation limit for an **RCID** and **AT**, the controller provides the **READ\_LIMIT** operation which can be requested by writing to the **cc\_alloc\_ctl** register. Upon successful completion of the operation, the **cc\_block\_mask** register holds the configured capacity allocation limit.

# Chapter 4. Bandwidth-controller QoS Register Interface

Controllers, such as memory controllers, that support bandwidth allocation and bandwidth usage monitoring provide a memory-mapped bandwidth-controller QoS register interface.

Table 8. Bandwidth-controller QoS register layout

Offset	Name	Size	Description	Optional?
0	<code>bc_capabilities</code>	8	Capabilities	No
8	<code>bc_mon_ctl</code>	8	Usage monitoring control	Yes
16	<code>bc_mon_ctr_val</code>	8	Monitoring counter value	Yes
24	<code>bc_alloc_ctl</code>	8	Bandwidth allocation control	Yes
32	<code>bc_bw_alloc</code>	8	Bandwidth allocation	Yes

The reset value is 0 for the following registers fields.

- `bc_mon_ctl` - `BUSY` field
- `bc_alloc_ctl` - `BUSY` field

The reset value is `UNSPECIFIED` for all other registers and/or fields.

The bandwidth controllers at reset must allocate all available bandwidth to `RCID` value of 0. When the bandwidth controller supports bandwidth allocation per access-type, the access-type value of 0 of `RCID=0` is allocated all available bandwidth, while all other access-types associated with that `RCID` share the bandwidth allocation with `AT=0`. The bandwidth allocation for all other `RCID` values is `UNSPECIFIED`. The bandwidth controller behavior in handling a request with a non-zero `RCID` value before configuring the bandwidth controller with bandwidth allocation for that `RCID` is also `UNSPECIFIED`.

## 4.1. Capabilities (`bc_capabilities`)

The `bc_capabilities` register is a read-only register that holds the bandwidth-controller capabilities.




Figure 6. Capabilities register fields

The `VER` field holds the version of the specification implemented by the bandwidth controller. The low nibble is used to hold the minor version of the specification and the upper nibble is used to hold the major version of the specification. For example, an implementation that supports version 1.0 of the specification reports 0x10.

The `NBWLKS` field holds the total number of available bandwidth blocks in the controller. The

bandwidth represented by each bandwidth block is **UNSPECIFIED**. The bandwidth controller supports reserving bandwidth in multiples of a bandwidth block, which enables proportional allocation of bandwidth.

Bandwidth controllers may limit the maximum bandwidth that may be reserved to a value smaller than **NBWBKLS**. The **MRBWB** field reports the maximum number of bandwidth blocks that can be reserved.



The bandwidth controller needs to meter the bandwidth usage by a workload to determine if it is exceeding its allocations and, if necessary, take necessary measures to throttle the workload's bandwidth usage. Therefore, the instantaneous bandwidth used by a workload may either exceed or fall short of the configured allocation. QoS capabilities are statistical in nature and are typically designed to enforce the configured bandwidth over larger time windows. By not allowing all available bandwidth blocks to be reserved for allocation, the bandwidth controller can handle such transient inaccuracies.

## 4.2. Bandwidth usage monitoring control (**bc\_mon\_ctl**)

The **bc\_mon\_ctl** register is used to control monitoring of bandwidth usage by a **MCID**. When the controller does not support bandwidth usage monitoring, the **bc\_mon\_ctl** register is read-only and hardwired to zero.




Figure 7. Bandwidth usage monitoring control (**bc\_mon\_ctl**)

Bandwidth controllers that support bandwidth usage monitoring implement a usage monitoring counter for each supported **MCID**. The usage monitoring counter may be configured to count a monitoring event. When an event matching the event configured for the **MCID** occurs then the monitoring counter is updated. The event matching may optionally be filtered by the access-type. The monitoring counter for bandwidth usage counts the number of bytes transferred by requests matching the monitoring event as the requests go past the monitoring point.

The **OP**, **AT**, **MCID**, and **EVT\_ID** fields of the register are WARL fields.

The **OP** field is used to instruct the controller to perform an operation listed in [Table 9](#). The supported operations of the **OP** field include configuring an event to count in the monitoring counter and obtaining a snapshot of the counter value.

The **EVT\_ID** field is used to program the identifier of the event to count in the monitoring counter selected by **MCID**. The **AT** field is used to program the access-type to count, and its validity is indicated by the **ATV** field. When **ATV** is 0, the counter counts requests with all access-types, and the **AT** value is ignored.

When the **EVT\_ID** for a **MCID** is programmed with a non-zero and legal value, the counter is reset to 0 and starts counting matching events for requests with the matching **MCID** and **AT** (if **ATV** is 1).

However, if the **EVT\_ID** is configured as 0, the counter retains its current value but stops counting.

A controller that does not support monitoring by access-type can hardwire the **ATV** and the **AT** fields to 0, indicating that the counter counts requests with all access-types.

Table 9. Usage monitoring operations (**OP**)

Operation	Encoding	Description
—	0	Reserved for future standard use.
<b>CONFIG_EVENT</b>	1	Configure the counter selected by <b>MCID</b> to count the event selected by <b>EVT_ID</b> , <b>AT</b> , and <b>ATV</b> . The <b>EVT_ID</b> encodings are listed in Table 10.
<b>READ_COUNTER</b>	2	Snapshot the value of the counter selected by <b>MCID</b> into <b>bc_mon_ctr_val</b> register. The <b>EVT_ID</b> , <b>AT</b> , and <b>ATV</b> fields are not used by this operation.
—	3-23	Reserved for future standard use.
—	24-31	Designated for custom use.

Table 10. Bandwidth monitoring event ID (**EVT\_ID**)

Event ID	Encoding	Description
<b>None</b>	0	Counter does not count and retains its value.
<b>Total Read and Write byte count</b>	1	Counter is incremented by the number of bytes transferred by a matching read or write request as the requests go past the monitor.
<b>Total Read byte count</b>	2	Counter is incremented by the number of bytes transferred by a matching read request as the requests go past the monitor.
<b>Total Write byte count</b>	3	Counter is incremented by the number of bytes transferred by a matching write request as the requests go past the monitor.
—	4-127	Reserved for future standard use.
—	128-256	Designated for custom use.

When the **bc\_mon\_ctl** register is written, the controller may need to perform several actions that may not complete synchronously with the write. A write to the **bc\_mon\_ctl** sets the read-only **BUSY** bit to 1 indicating the controller is performing the requested operation. When the **BUSY** bit reads 0, the operation is complete, and the read-only **STATUS** field provides a status value (see Table 11 for details). Written values to the **BUSY** and the **STATUS** fields are ignored. An implementation that can complete the operation synchronously with the write may hardwire the **BUSY** bit to 0. The state of the **BUSY** bit, when not hardwired to 0, shall only change in response to a write to the register. The **STATUS** field remains valid until a subsequent write to the **bc\_mon\_ctl** register.

Table 11. **bc\_mon\_ctl.STATUS** field encodings

<b>STATUS</b>	Description
0	Reserved

STATUS	Description
1	The operation was successfully completed.
2	An invalid operation ( <b>OP</b> ) was requested.
3	An operation was requested for an invalid <b>MCID</b> .
4	An operation was requested for an invalid <b>EVT_ID</b> .
5	An operation was requested for an invalid <b>AT</b> .
6-63	Reserved for future standard use.
64-127	Designated for custom use.

When the **BUSY** bit is set to 1, the behavior of writes to the **bc\_mon\_ctl** is **UNSPECIFIED**. Some implementations may ignore the second write, while others may perform the operation determined by the second write. To ensure proper operation, software must first verify that the **BUSY** bit is 0 before writing the **bc\_mon\_ctl** register.

### 4.3. Bandwidth monitoring counter value (**bc\_mon\_ctr\_val**)

The **bc\_mon\_ctr\_val** is a read-only register that holds a snapshot of the counter selected by a **READ\_COUNTER** operation. When the controller does not support bandwidth usage monitoring, the **bc\_mon\_ctr\_val** register always reads as zero.




Figure 8. Bandwidth monitoring counter value (**bc\_mon\_ctr\_val**)

The counter is valid if the **INV** field is 0. The counter may be marked **INV** if, for **UNSPECIFIED** reasons, the controller determine the count to be not valid. Such counters may become valid in the future. Additionally, if an unsigned integer overflow of the counter occurs, then the **OVF** bit is set to 1.



A counter may be marked as **INV** if the controller has not been able to establish an accurate counter value for the monitored event.

The counter provides the number of bytes transferred by requests matching the **EVT\_ID** as they go past the monitoring point. A bandwidth value may be determined by reading the byte count value at two instances of time **T1** and **T2**, see [equation \(2\)](#). If the value of the counter at time **T1** was **B1**, and at time **T2** is **B2**, then the bandwidth can be calculated as follows. The frequency of the time source is represented by  $T_{freq}$ .

$$Bandwidth = T_{freq} \times \frac{B2 - B1}{T2 - T1} \quad (2)$$

The width of the counter is **UNSPECIFIED**.



While the width of the counter is **UNSPECIFIED**, it is recommended to be wide enough to prevent more than one overflow per sample when the sampling frequency is 1 Hz.

If an overflow was detected then software may discard that sample and reset the counter and overflow indication by reprogramming the event using **CONFIG\_EVENT** operation.

## 4.4. Bandwidth allocation control (**bc\_alloc\_ctl**)

The **bc\_alloc\_ctl** register is used to control the allocation of bandwidth to an **RCID** per **AT**. If a controller does not support bandwidth allocation, then the register is read-only zero. If the controller does not support bandwidth allocation per access-type, then the **AT** field is also read-only zero.




Figure 9. Bandwidth allocation control (**bc\_alloc\_ctl**)

The **OP** field instructs the bandwidth controller to perform an operation listed in Table 12. The **bc\_alloc\_ctl** register is used in conjunction with the **bc\_bw\_alloc** register to perform bandwidth allocation operations. If the requested operation uses the operands configured in **bc\_bw\_alloc**, software must first program the **bc\_bw\_alloc** register with the operands for the operation before requesting the operation.

Table 12. Bandwidth allocation operations (**OP**)

Operation	Encoding	Description
—	0	Reserved for future standard use.
CONFIG_LIMIT	1	Establishes reserved bandwidth allocation for requests by <b>RCID</b> and of access-type <b>AT</b> . The bandwidth allocation is specified in <b>bc_bw_alloc</b> register.
READ_LIMIT	2	Reads back the previously configured bandwidth allocation for requests by <b>RCID</b> and of access-type <b>AT</b> . The current configured allocation is written to <b>bc_bw_alloc</b> register on completion of the operation.
—	3-23	Reserved for future standard use.
—	24-31	Designated for custom use.

A bandwidth allocation must be configured for each access-type supported by the controller. When differentiated bandwidth allocation based on access-type is not required, one of the access-types may be designated to hold a default bandwidth allocation, and the other access-types can be configured to share the allocation with the default access-type. If bandwidth is not allocated for each access-type supported by the controller, the behavior is **UNSPECIFIED**.

When the `bc_alloc_ctl` register is written, the controller may need to perform several actions that may not complete synchronously with the write. A write to the `bc_alloc_ctl` sets the read-only `BUSY` bit to 1 indicating the controller is performing the requested operation. When the `BUSY` bit reads 0, the operation is complete, and the read-only `STATUS` field provides a status value (see Table 13 for details). Written values to the `BUSY` and the `STATUS` fields are ignored. An implementation that can complete the operation synchronously with the write may hardwire the `BUSY` bit to 0. The state of the `BUSY` bit, when not hardwired to 0, shall only change in response to a write to the register. The `STATUS` field remains valid until a subsequent write to the `bc_alloc_ctl` register.

Table 13. `bc_alloc_ctl.STATUS` field encodings

STATUS	Description
0	Reserved
1	The operation was successfully completed.
2	An invalid operation ( <code>OP</code> ) was requested.
3	An operation was requested for an invalid <code>RCID</code> .
4	An operation was requested for an invalid <code>AT</code> .
5	An invalid or unsupported reserved bandwidth block was requested.
6-63	Reserved for future standard use.
64-127	Designated for custom use.

## 4.5. Bandwidth allocation configuration (`bc_bw_alloc`)

The `bc_bw_alloc` is used to program reserved bandwidth blocks (`Rbwb`) for an `RCID` for requests of access-type `AT` using the `CONFIG_LIMIT` operation. If a controller does not support bandwidth allocation, then the `bc_bw_alloc` register is read-only zero.

The `bc_bw_alloc` holds the previously configured reserved bandwidth blocks for an `RCID` and `AT` on successful completion of the `READ_LIMIT` operation.

Bandwidth is allocated in multiples of bandwidth blocks, and the value in `Rbwb` must be at least 1 and must not exceed `MRBWB`. Otherwise, the `CONFIG_LIMIT` operation fails with `STATUS=5`. Additionally, the sum of `Rbwb` allocated across all `RCID`'s must not exceed `MRBWB`, or the `CONFIG_LIMIT` operation fails with `STATUS=5`.




Figure 10. Bandwidth allocation configuration (`bc_bw_alloc`)

The `Rbwb`, `Mweight`, `sharedAT`, and `useShared` are all WARL fields.



Bandwidth allocation is typically enforced by the bandwidth controller over finite accounting windows. The process involves measuring the bandwidth consumption over an accounting window and using the measured bandwidth to determine if an **RCID** is exceeding its bandwidth allocations for each access-types. The specifics of how the accounting window is implemented are **UNSPECIFIED**, but is expected to provide a statistically accurate control of the bandwidth usage over a few accounting intervals.

The **Rbwb** represents the bandwidth that is made available to a **RCID** for requests matching **AT**, even when all other **RCID** are using their full allocation of bandwidth.

If there is non-reserved or unused bandwidth available in an accounting interval, **RCID**'s may compete for additional bandwidth. The non-reserved or unused bandwidth is proportionately shared among the competing **RCID**s using the configured **Mweight** parameter, which is a number between 0 and 255. A larger weight implies a greater fraction of the bandwidth. A weight of 0 implies that the configured limit is a hard limit, and the use of unused or non-reserved bandwidth is not allowed.

The sharing of non-reserved bandwidth is not differentiated by access-type. Therefore, the **Mweight** parameter must be programmed identically for all access-types. If this parameter is programmed differently for each access-type, then the controller may use the parameter configured for any of the access-types, but the behavior is otherwise well defined.

When the **Mweight** parameter is not set to 0, the amount of unused bandwidth allocated to **RCID=x** during contention with another **RCID** that is also permitted to use unused bandwidth is determined by dividing the **Mweight** of **RCID=x** by the sum of the **Mweight** of all other contending **RCID**'s. This ratio **P** is determined as in [equation \(3\)](#).

$$P = \frac{Mweight_x}{\sum_{r=1}^n Mweight_r} \quad (3)$$



The bandwidth enforcement is typically work-conserving, meaning that it allows unused bandwidth to be used by requestors enabled to use it even if they have consumed their **Rbwb**.

When contending for unused bandwidth, the weighted share is typically computed among the **RCID**'s that are actively generating requests in that accounting interval and have a non-zero weight programmed.

If unique bandwidth allocation is not required for an access-type, then the **useShared** parameter may be set to 1 for a **CONFIG\_LIMIT** operation. When **useShared** is set to 1, the **sharedAT** field specifies the access-type with which the bandwidth allocation is shared by the access-type in **bc\_alloc\_ctl.AT**. In this case, the **Rbwb** and **Mweight** fields are ignored, and the configurations of the access-type in **sharedAT** are applied. If the access-type specified by **sharedAT** does not have unique bandwidth allocation, meaning that it has not been configured with **useShared=0**, then the behavior is **UNSPECIFIED**.

The **useShared** and **sharedAT** fields are reserved if the bandwidth controller does not support bandwidth allocation per access-type.



When unique bandwidth allocation for an access-type is not required then one or more access-types may be configured with a shared bandwidth allocation. For example, consider a bandwidth controller that supports 3 access-types. The access-type 0 and 1 of **RCID** 3 are configured with unique bandwidth allocations and the access-type 2 is configured to share bandwidth allocation with access-type 1. The example configuration is illustrated as follows:

	Rbwb	Mweight	useShared	sharedAT
RCID=3, AT=0	100	16	0	N/A
RCID=3, AT=1	50	16	0	N/A
RCID=3, AT=2	N/A	N/A	1	1

# Chapter 5. IOMMU extension for QoS ID

Monitoring or allocation of resources accessed by the IOMMU and devices governed by the IOMMU requires a way to associate QoS IDs with such requests. This section specifies a RISC-V IOMMU [3] extension to:

- Configure and associate QoS IDs for device-originated requests.
- Configure and associate QoS IDs for IOMMU-originated requests.

## 5.1. IOMMU registers

The specified memory-mapped register layout defines a new IOMMU register named `iommu_qosid`. This register is used to configure the Quality of Service (QoS) IDs associated with IOMMU-originated requests. The register has a size of 4 bytes and is located at an offset of 624 from the beginning of the memory-mapped region.

Table 14. IOMMU Memory-mapped register layout

Offset	Name	Size	Description	Is Optional?
624	<code>iommu_qosid</code>	4	QoS IDs for IOMMU requests.	Yes
628	Reserved	60	Reserved for future use (WPRI)	

### 5.1.1. Reset behavior

If the reset value for `ddtp.iommu_mode` field is `Bare`, then the `iommu_qosid.RCID` field must have a reset value of 0.



At reset, it is suggested that the `RCID` field of `iommu_qosid` be set to 0 if the IOMMU is in `Bare` mode, as typically the resource controllers in the SoC default to a reset behavior of associating all capacity or bandwidth to the `RCID` value of 0.

### 5.1.2. IOMMU capabilities (`capabilities`)

The IOMMU capabilities register has been extended with a new field, `QOSID`, which enumerate support for associating QoS IDs with requests made through the IOMMU.



## 5.2. Device-context fields

The **ta** field of the device context is extended with two new fields, **RCID** and **MCID**, to configure the QoS IDs to associate with requests originated by the devices.




Figure 13. Translation attributes (**ta**) field

The **RCID** and **MCID** configured in **DC.ta** are associated with IOMMU-initiated requests to access the following data structures:

- Process directory table (**PDT**)
- Second-stage page table
- First-stage page table
- MSI page table
- Memory-resident interrupt file (**MRIF**)

The **RCID** and **MCID** configured in **DC.ta** are provided to the IO bridge on successful address translations. The IO bridge should associate these QoS IDs with device-initiated requests.

If **capabilities.QOSID** is 1 and **DC.ta.RCID** or **DC.ta.MCID** is wider than that supported by the IOMMU, a **DC** with **DC.tc.V=1** is considered misconfigured. In this case, the IOMMU should stop and report "DDT entry misconfigured" (cause = 259).

## 5.3. IOMMU ATC capacity allocation and monitoring

The IOMMU may support capacity allocation and usage monitoring in the IOMMU address translation cache (IOATC) by implementing a capacity controller register interface.

Some IOMMU may support multiple IOATC where the capacity of each such IOATC may not be the same (e.g., corresponding each page sizes supported). When multiple IOATC are implemented, the IOMMU may implement a capacity controller register interface for each IOATC to enable capacity allocation in each IOATC.

# Chapter 6. Hardware Guidelines

## 6.1. Sizing QoS Identifiers

In a typical implementation the number of **RCID** bits implemented (e.g., to support 10s of **RCIDs**) may be smaller than the number of **MCID** bits implemented (e.g., to support 100s of **MCIDs**).

It is a typical usage to associate a group of applications/VMs with a common **RCID** and thus sharing a common pool of resource allocations. The resource allocations for the **RCID** is established to meet the SLA objectives of all members of the group. If SLA objectives of one or more members of the group stop being met, the resource usage of one or more members of the group may be monitored by associating them with a unique **MCID** and this iterative analysis process used to determine the optimal strategy - increasing resources allocated to the **RCID**, moving some members to a different **RCID**, migrating some members away to another machine, etc. - for restoring the SLA. Having a sufficiently large pool of **MCID** speeds up this analysis.



To support maximal flexibility in allocation of QoS IDs to workloads it is recommended for all resource controllers in the system to support an identical number of **RCID** and **MCID**.

## 6.2. Sizing monitoring counters

Typically software samples the monitoring counters periodically to monitor capacity and bandwidth usage. The width of the monitoring counters is recommended to be wide enough to not cause more than one overflow per sample when sampled at a frequency of 1 Hz.

# Chapter 7. Software Guidelines

## 7.1. Reporting capacity and bandwidth controllers

The capability and bandwidth controllers that are present in the system should be reported to operating systems using methods such as ACPI and/or device tree. For each capacity and bandwidth controller, the following information should be reported using these methods:

- Type of controller (e.g, cache, interconnect, memory, etc.)
- Location of the register programming interface for the controller
- Placement and topology describing the hart and IO bridges that share the resources controlled by the controller
- The number of QoS identifiers supported by the controller

## 7.2. Context switching QoS Identifiers

Typically, the contents of the `sqoscfg` CSR are updated with a new `RCID` and/or `MCID` by the HS/S-mode scheduler if the `RCID` and/or `MCID` of the new workload (a process or a VM) is not same as that of the old workload.

A context switch usually involves saving the context associated with the workload being switched away from and restoring the context of the workload being switched to. Such context switch may be invoked in response to an explicit call from the workload (i.e, as a function of an `ECALL` invocation) or may be done asynchronously (e.g., in response to a timer interrupt). In such cases the scheduler may want to execute with the `sqoscfg` configuration of the workload being switched away from such that this execution is attributed to the workload being switched away from and then prior to restoring the new workloads context, first switch to the `sqoscfg` configuration appropriate for the workload being switched to such that all of that execution is attributed to the new workload. Further in this context switch process, if the scheduler intends some of its execution to be attributed to neither the outgoing workload nor the incoming workload, then the scheduler may switch to a new `sqoscfg` configuration that is different from that of either of the workloads for the duration of such execution.

## 7.3. QoS configurations for virtual machines

Usually for virtual machines the resource allocations are configured by the hypervisor. Usually the Guest OS in a virtual machine does not participate in the QoS flows as the Guest OS does not know the physical capabilities of the platform or the resource allocations for other virtual machines in the system.

If a use case requires it, a hypervisor may virtualize the QoS capability to a VM by virtualizing the memory-mapped CBQRI register interface and virtualizing the virtual-instruction exception on access to `sqoscfg` CSR by the Guest OS.



If the use of directly selecting among a set of `RCID` and/or `MCID` by a VM becomes

more prevalent and the overhead of virtualizing the `sqoscfg` CSR using the virtual instruction exception is not acceptable then a future extension may be introduced where the `RCID/MCID` attempted to be written by VS mode are used as a selector for a set of `RCID/MCID` that the hypervisor configures in a set of HS mode CSRs.

## 7.4. QoS Identifiers for supervisor and machine mode

The `RCID` and `MCID` configured in `sqoscfg` also apply to execution in S/HS-mode, but this is typically not an issue. Usually, S/HS-mode execution occurs to provide services, such as through the SBI, to software executing at lower privilege. Since the S/HS-mode invocation was to provide a service for the lower privilege mode, the S/HS-mode software may not modify the `sqoscfg` CSR.

Similarly, The `RCID` and `MCID` configured in `sqoscfg` also apply to execution in M-mode, but this is typically not an issue either. Usually, M-mode execution occurs to provide services, such as through the SBI interface, to software executing at lower privilege. Since the M-mode invocation was to provide a service for the lower privilege mode, the M-mode software may not modify the `sqoscfg` CSR.

If separate `RCID` and/or `MCID` are needed during software execution in M/S/HS-mode, then the M/S/HS-mode software may update the `sqoscfg` CSR and restore it before returning to lower privilege mode execution. The statistical nature of QoS capabilities means that the brief duration, such as the few instructions in the M/S/HS-mode trap handler entry point, during which the trap handler may execute with the `RCID` and/or `MCID` established for lower privilege mode operation may not have a significant statistical impact.



# Bibliography

- [1] K. Du Bois, S. Eyerman, and L. Eeckhout, “Per-Thread Cycle Accounting in Multicore Processors,” *ACM Trans. Archit. Code Optim.*, vol. 9, no. 4, Jan. 2013, doi: 10.1145/2400682.2400688.
- [2] D. Lo, L. Cheng, R. Govindaraju, P. Ranganathan, and C. Kozyrakis, “Heracles: Improving Resource Efficiency at Scale,” in *Proceedings of the 42nd Annual International Symposium on Computer Architecture*, New York, NY, USA, 2015, pp. 450–462, doi: 10.1145/2749469.2749475.
- [3] “RISC-V IOMMU Architecture Specification.” [Online]. Available: [github.com/riscv/riscv-iommu](https://github.com/riscv/riscv-iommu).
- [4] N. C. Thornock and J. K. Flanagan, “Facilitating level three cache studies using set sampling,” in *2000 Winter Simulation Conference Proceedings (Cat. No.00CH37165)*, 2000, vol. 1, pp. 471–479 vol.1, doi: 10.1109/WSC.2000.899754.